AI PRIME STUDIO LUMINA

Template Prompt Operating Standard v1.2

Owner reference for future Photo / Video / Couple template prompts
Date: 04 November 2026
Status: Addendum to Prompt Standard v1.1 — does not replace the existing Master Contract or runtime Global Prompts.

1. Core rule

Create a short, exact, template-specific Scene Recipe on top of the existing protected Global Prompt branches. Do not paste one giant universal prompt into every template. The assistant must first understand the source media, desired edit, customer inputs, fixed template assets and the provider's actual capabilities.

The prompt text describes the visual task. Model, provider, resolution, duration, price, customer credits and storage are configuration/runtime data, not prompt prose.

2. Choose the correct prompt branch

Template type	Runtime prompt branch	Required separation
Standard Photo	PHOTO_GLOBAL_PROMPT + Photo Recipe	Customer photo supplies identity; designated template AFTER reference supplies target scene/pose/composition.
Standard Video	VIDEO_GLOBAL_PROMPT + selected Video Workflow + Video Recipe	Source video supplies source scene/performance/motion; each customer reference supplies only its assigned subject identity.
Couple Photo	Photo branch + COUPLE_PHOTO_MAPPING + Couple Photo Recipe	Person A and Person B are separate configured slots and identities.
Couple Video	Video branch + verified workflow + COUPLE_VIDEO_MAPPING + Couple Video Recipe	Only enabled after the selected provider is proven to preserve two identities in the real workflow.

Never confuse output type with upload type. template.type means the result is Photo or Video. requiredInputs / inputType determine what the customer uploads. A Video output may require a Photo input. A Photo output must not inherit a Video prompt.

Couple is a category/identity mode, not a replacement for modality. Couple Photo and Couple Video need separate prompt branches.

3. Assistant's process whenever a new template is sent

1. Inspect the supplied photo/video and any reference screenshots. Record duration, aspect ratio, visible subjects, scene changes/cutaways, camera movement, visual style and audio only when these are actually observable.
2. Clarify the exact transformation: what changes, what must remain, and whether the task is a local element replacement, source-preserving edit, motion transfer, photo edit or new video generation.
3. Assign every source asset a role: source video, customer identity, Person A, Person B, product/object reference, template-owned fixed asset, poster/cover or preview. Never treat a poster as an identity reference.
4. Build an exact subject/shot map. State which person or element is present in which scene. Do not invent a presence, identity, clothing detail or missing asset.
5. Create a concise English Scene Recipe. Reuse the existing Global Prompt; do not duplicate its general identity, realism, wardrobe or audio rules.
6. Give the Admin setup: output type, workflow, customer input slots, template-owned assets, provider/model candidates, verified parameters, estimated Credits basis and tests still required.
7. Label every uncertain capability, price or duration UNKNOWN / PENDING_QUOTE / NOT_VERIFIED; never present assumptions as facts.

4. Scene Recipe structure (English, Admin-ready)

Use only the sections needed by that template, in this order:

1. **TASK / CHANGE SCOPE** — exactly what is being changed and what is not being changed.
2. **SOURCE AND REFERENCE ROLES** — what the source video/photo controls; what each reference image controls.
3. **SUBJECT / SLOT MAP** — stable slot IDs and named subjects (Person A, Person B, Subject C, Product A). Do not rely only on image-array position.
4. **SHOT / SCENE MAP** — for multi-shot videos, map subjects to the scenes where they actually appear; note cutaways and exceptions.
5. **IDENTITY / APPEARANCE** — preserve recognizable identity and listed visible traits; do not redesign or blend separate identities.
6. **WARDROBE SOURCE** — Customer Photo, Template/Source, or Auto. Auto uses customer clothing only when clearly visible; otherwise use the designated template/source clothing. Never invent unseen clothing.
7. **PRESERVATION RULES** — list the non-target scene properties to preserve, but only when the selected workflow/model can reasonably preserve them.
8. **REALISM / CONTINUITY** — natural perspective, lighting integration and temporal consistency; avoid flicker, identity drift and visual artifacts.
9. **AUDIO** — request source-audio preservation only when the exact workflow supports it.
10. **DO-NOT-CHANGE LIST** — concise exclusions, no new scenes/subjects or unrelated additions unless explicitly requested.

Do not repeat the same instruction in every section. Resolve conflicts by priority: explicit change scope > assigned identity/asset roles > source video or template reference > optional customer instruction. Optional customer text cannot silently override protected template rules.

5. Customer uploads versus template-owned assets

- Customer-upload slots are for values that genuinely vary by customer (for example, Customer Person A's photo).
- Fixed scene elements belong to the Admin's template assets or source video, not to a compulsory customer upload slot.
- If an object/product is already clearly present in the source video, that video may be its reference only when the selected model's workflow supports preserving/using it.
- If a stable separate reference is needed, the Admin may select an approved still/reference from the source during template setup and save it under a separate template_asset role. Do this at setup time, not by silently extracting random frames on every paid generation.
- If the required item is not present in the customer input or approved template assets, stop and ask for the missing asset. Do not hallucinate or fetch an unrelated Gallery item.
- Reuse the existing slot/recipe schema where possible. If a field is genuinely missing, use structured slotId, role, media type, source (CUSTOMER_UPLOAD / TEMPLATE_ASSET / optional), required flag and ordering. Do not rebuild the input system.

6. Workflow and model selection

Select by task capability first, verified cost second:

Task	Preferred candidate	Important distinction
Local replacement of a selected subject/object while keeping other scene elements	Higgsfield Genjutsu Light Object Swap	Candidate for selected-element replacement; verify exact API endpoint, request schema and account support.
Transfer source movement/performance to a new subject or scene	Higgsfield Genjutsu Light Motion Transfer	Not equivalent to exact local replacement; scene reconstruction may differ.
Standard Photo / Couple Photo	Existing Google/Gemini image adapter candidate	Verify image output, multiple-reference handling, price and privacy before production.
General text/image/video-reference generation or compatible video edit	Cheapest verified compatible model from the current provider catalogue	Do not auto-enable a model solely because it appears in the UI/catalogue. Verify cost, privacy, commercial terms and actual output.

Task	Preferred candidate	Important distinction
Gemini Omni	Evaluation-only until approved test gates pass	Never use as a hidden production fallback.

Admin must retain Auto / Manual choice. Auto chooses the lowest-cost model only from the verified-compatible allowlist for that template's workflow, input roles, output type, duration and quality. Manual shows only models compatible with that same contract. Higher-cost fallback needs a new quote and customer consent. No silent provider/model switch.

7. Quality, duration and Credits

- Default quality is the lowest quality verified and supported by that exact model endpoint/account. A general 480p assumption must not be applied to all providers.
- A 5-second default and 8-second quick option are sensible only when the selected endpoint truly accepts/controls requested output duration. Show a duration slider only for supported values.
- Some video-edit modes ignore a duration input and use the source clip's duration. For such a model, hide/disable the 5s/8s output-duration selector and price using the verified source-duration billing rule.
- Genjutsu Light's current official API playground lists 480p / 720p / 1080p and per-second list prices of $0.05 / $0.10 / $0.20 respectively. It states billable seconds are output duration plus the sum of input reference-video durations, input durations are measured during preprocessing, the source is trimmed to the first 15 seconds, and output duration is rounded up to whole seconds. These are published list prices, not an account-specific quote. Re-check the exact endpoint and account before enabling a customer quote.
- Keep separate: published list price, account quote/discount, estimated provider cost, actual provider transaction cost, estimated customer Credits and final reconciled Credits.
- Required Credits formula: ceil(verifiedProviderCostUSD × configuredUSDINRRate × 1.40 × creditsPerINR). Apply the 40% markup exactly once.
- If model capability, duration behavior, billable-duration rule or price is unknown, return PENDING_QUOTE and block charge/generation. Never label a list-price estimate as actual usage cost. The customer sees Credits only, never INR/USD on the Generation screen.

8. Fast media loading

Use this sequence: poster immediately → lightweight preview near viewport → background HD prefetch when safe → seamless HD playback upgrade.

- Preserve the existing poster-first and lazy-preview system; do not attach every full video URL on initial page load.
- Schedule background HD after visible/near-visible priority and browser idle. Priority: clicked item > visible item > remaining gallery items.
- On Data Saver or a slow connection, suppress automatic HD/background prefetch and retain the lightweight preview. Do not force a large download.
- Upgrade the current player without resetting currentTime, paused/playing state or mute state. Avoid flicker/restart.
- Limit concurrency and cache size; revoke unused Blob URLs; use progressive/CDN range delivery where supported. Do not download every HD file at once.
- Keep the original generation/source video unchanged. Poster, preview and HD derivative are separate media roles.
- Do not add HLS/transcoding infrastructure until traffic and cost justify it.

9. Higgsfield app/API and asset rights

Higgsfield's API catalog and its website UI/preset library are not automatically the same thing. Only documented API models, endpoints, parameters and preset calls can be integrated. Do not scrape, copy, re-host or resell Higgsfield's proprietary templates, preset assets or sample media as a standalone library without explicit rights. Use owner-created, customer-authorized or properly licensed assets.

The current Terms say the user is responsible for rights in uploaded content; the standard Terms also describe processing and model-improvement use of inputs/outputs. Before private customer photos are sent to a provider, verify the applicable privacy/training terms for the actual account or agreement.

10. Required answer format from the Assistant for every new template

1. Final recommendation — workflow and why.
2. Admin-ready English Scene Recipe — ready to paste into the one template's recipe field.
3. Admin configuration — output type, required customer slots, template-owned assets, wardrobe option, provider/model candidate, verified quality/duration controls.
4. Cost/credit basis — only when verified; otherwise clearly PENDING_QUOTE.
5. Acceptance checklist — preview, test, input mapping, output, playback, cost and fallback checks.

Do not return only a generic prompt. The Scene Recipe must map to the actual template's assets and configured slots. Do not modify the application's Global Prompt or unrelated settings just to accommodate one template.

Official reference links (re-check before production changes)

- Genjutsu Light Motion Transfer API Playground: https://open.higgsfield.ai/models/higgsfield/genjutsu-light/motion-transfer/v1.0/playground
- Genjutsu Light Object Swap API Playground: https://open.higgsfield.ai/models/higgsfield/genjutsu-light/object-swap/v1.0/playground
- Higgsfield Terms of Use: https://higgsfield.ai/terms-of-use-agreement
- Higgsfield API overview/help: https://higgsfield.ai/creator-hub/help-center/integrations/what-is-the-higgsfield-api